

k8s complete setup -master & worker node

properly, from scratch, clean DevOps-style 💪

I'll assume you want a **basic Kubernetes cluster** using **kubeadm** on **Linux (Ubuntu 20.04/22.04)** with:

- 🧠 Master node → `10.0.2.14`
- 🏠 Worker node → `10.0.2.190`
- Container runtime: **containerd**
- Network plugin: **Calico** (industry-standard)

I'll give you **end-to-end steps**, copy-paste friendly, with **why** each step exists.

Architecture Overview

Master Node (Control Plane)

IP: `10.0.2.14`

```
├── kube-apiserver  
├── etcd  
├── controller-manager  
└── scheduler
```

Worker Node

IP: `10.0.2.190`

```
└── kubelet + kube-proxy
```

◆ STEP 0: Prerequisites (BOTH NODES)

OS & User

- Ubuntu 20.04 / 22.04
- Login as **root** or a sudo user

Set hostnames

Master

```
hostnamectl set-hostname k8s-master
```

Worker

```
hostnamectl set-hostname k8s-worker
```

Logout & login again.

Update system

```
sudo apt update && sudo apt upgrade -y
```

Disable swap (VERY IMPORTANT)

```
sudo swapoff -a  
sudo sed -i '/ swap / s/^/#/' /etc/fstab
```

Verify:

```
free -h
```

Enable kernel modules

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf  
overlay  
br_netfilter  
EOF
```

```
sudo modprobe overlay  
sudo modprobe br_netfilter
```

Sysctl settings (networking)

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf  
net.bridge.bridge-nf-call-iptables = 1  
net.bridge.bridge-nf-call-ip6tables = 1  
net.ipv4.ip_forward = 1  
EOF
```

```
sudo sysctl --system
```

◆ STEP 1: Install Container Runtime (containerd) – BOTH NODES

```
sudo apt install -y containerd
```

Configure containerd

```
sudo mkdir -p /etc/containerd  
sudo containerd config default | sudo tee /etc/containerd/config.toml
```

Edit config:

```
sudo nano /etc/containerd/config.toml
```

Find:

```
SystemdCgroup = false
```

Change to:

```
SystemdCgroup = true
```

Restart:

```
sudo systemctl restart containerd  
sudo systemctl enable containerd
```

Check:

```
containerd --version
```

◆ STEP 2: Install Kubernetes Components – BOTH NODES

Add Kubernetes repo

```
sudo apt install -y apt-transport-https ca-certificates curl  
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --  
dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg  
echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]  
https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /" | sudo tee  
/etc/apt/sources.list.d/kubernetes.list
```

Update & install:

```
sudo apt update  
sudo apt install -y kubelet kubeadm kubectl
```

Hold versions:

```
sudo apt-mark hold kubelet kubeadm kubectl
```

◆ STEP 3: Initialize MASTER Node (10.0.2.14)

⚠ Run ONLY on master

```
sudo kubeadm init \
--apiserver-advertise-address=10.0.2.14 \
--pod-network-cidr=192.168.0.0/16
```

 Wait 2–3 minutes.

Configure kubectl (MASTER)

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Test:

```
kubectl get nodes
```

You'll see:

```
k8s-master    NotReady
```

 This is normal (network plugin not installed yet)

◆ STEP 4: Install Network Plugin (Calico) – MASTER

```
kubectl apply -f
https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/calico.yaml
```

Wait:

```
kubectl get pods -n kube-system
```

After 1–2 mins:

```
calico-node    Running
```

Now:

```
kubectl get nodes
```

Master should be:

```
Ready
```

◆ STEP 5: Join WORKER Node (10.0.2.190)

During `kubeadm init`, you got a join command like this 👇

Run it on the worker:

```
sudo kubeadm join 10.0.2.14:6443 \
--token <token> \
--discovery-token-ca-cert-hash sha256:<hash>
```

If token expired, regenerate on master:

```
kubeadm token create --print-join-command
```

◆ STEP 6: Verify Cluster (MASTER)

```
kubectl get nodes -o wide
```

Expected:

```
k8s-master    Ready   control-plane   10.0.2.14
k8s-worker    Ready   <none>        10.0.2.190
```

🎉 Cluster is UP

◆ STEP 7: (Optional) Allow Pods on Master (Lab only)

```
kubectl taint nodes k8s-master node-role.kubernetes.io/control-plane-
```

⚠️ Do NOT do this in production.

◆ STEP 8: Test with Sample App

```
kubectl create deployment nginx --image=nginx
kubectl expose deployment nginx --port=80 --type=NodePort
kubectl get svc
```

🧠 Common Troubleshooting

Issue	Fix
Node NotReady	Check Calico
kubeadm init fails	Swap not disabled
Pods Pending	Network plugin missing

Issue	Fix
Worker not joining	Firewall / token expired

🏁 Resume-Ready Line

Built a multi-node Kubernetes cluster using kubeadm with containerd runtime, Calico networking, and manual control-plane & worker node configuration.

If you want next:

- 🏁 HA master setup
- 🔒 TLS / RBAC deep dive
- 🚀 Deploy Dockerized MERN app
- 🛠 Jenkins → Kubernetes deployment
- 📄 Diagram + YAML examples

Just say **next** 